

Logistic Regression

Training Time: 1629ms

weights = [0.999877, -2.41086]

Accuracy: 0.784553

Sensitivity: 0.695652

Specificity: 0.862595

Naïve Bayes:

Prior Probabilities: Yes=0.39, No=0.61

Accuracy: 0.784553

Sensitivity: 0.695652

Specificity: 0.862595

Interestingly, I would have expected the two algorithms to perform slightly differently to each other. But both algorithms had the exact same amount of true positives, true negatives, false positives, and false negatives.